

Life Before Language: Separating the Engineering Gap from the Metaphysical Gap in Debates on Machine Consciousness

Alex Bone

Independent Researcher, Long Beach, CA

Preprint - July 2026

Abstract

Contemporary debates about artificial intelligence and consciousness routinely commit a category error: they compare human and machine cognition at their endpoints (fluent language, symbolic reasoning) rather than at their origins. Human consciousness does not begin with language; it begins with a pre-rational developmental interface - appetite, attachment, sensation, consequence, and temporal continuity - through which reasoning later emerges as the act of a self. Large language models invert this ordering: they begin at the symbolic layer and possess no developmental substrate beneath it. This paper makes four contributions. First, it formalizes a seven-level hierarchy of the developmental interface and locates current AI architectures within it. Second, it distinguishes two gaps that are routinely conflated: an *engineering gap* (bridgeable in principle) and a *metaphysical gap* (possibly not). Third, it proposes and executes a two-workstream research program in which four implemented probe systems - an entropy-decay decoder, a spectral-geometry regulator, a holographic memory layer, and a bearer-indexed continuant agent - each realize one fragment of the developmental interface, and a flagship *construction-order experiment* introduces the same regulatory geometry at three points in a model's formation. Fourth, it reports the results. The construction-order experiment yields three stable regimes - inference-only overlays are a *tax*, fine-tuned retrofits are a *crutch*, and constraints present from the first training step are *internalized* (removable at almost no cost, their discipline persisting in the weights) - refining the framework's original modulation/answerability binary into a graded ladder: modulation -> internalization -> **the wall** -> consequence -> subjectivity. Retrofit adaptation approaches but does not reach the grown-in signature across all tested regimes (path dependence with asymptotic non-convergence), and bearer-identity in the continuant probe becomes load-bearing exactly when outcomes stop being facts about the world and become facts of the bearer. Nothing measured crosses the wall: the grown-in model demonstrates that constitution and consequence-bearing are separable - it was changed by its formation, and nothing mattered to it. The experiments locate the wall; they do not dissolve it.

Keywords: philosophy of mind, machine consciousness, developmental cognition, hallucination, embodiment, construction order, internalization, Leibniz, Boethius

1. Introduction: The Category Error

The question "can AI reason?" is asked as though reasoning were the natural starting point of mind. It is not. In human beings, reasoning is a late arrival - an achievement built on top of a pre-existing order of hunger, fear, attachment, imitation, and sensation. A newborn does not conclude "I am hungry, therefore I should request food." It cries. There is an inner regulatory order before language, before argument, before self-explanation.

Large language models begin at the opposite end. They are instantiated directly at the symbolic layer,

trained on the linguistic residue of beings who came to language through life. The comparison between human and machine intelligence therefore fails at the point of origin, not at the point of performance. The slogan form of the thesis is:

AI begins with language without life. Human consciousness begins with life before language.

This is not, by itself, an argument that machine consciousness is impossible. It is an argument that the debate has been conducted at the wrong layer. This paper reorganizes the debate around origins rather than endpoints, separates the engineering question from the metaphysical one, and - unlike the purely philosophical treatments it descends from - reports measurements. The framework's sharpest questions were initially posed as open; this paper answers several of them at small scale, with the honest bounds that scale imposes.

2. The Developmental Interface Hierarchy

We define the *developmental interface* as the total ordered arrangement through which a human consciousness enters, engages, and is shaped by the world. We stratify it into seven levels, ordered from ontological ground to symbolic surface:

Level 1 - Providential order. The total order of being, possibility, and final purpose; in Boethian terms, the whole seen from outside temporal sequence (Boethius, *Consolation of Philosophy*, Bk. IV-V).

Level 2 - Soul / primary consciousness. The living center of experience that enters time; the subject for whom there is something it is like (Nagel, 1974) to unfold.

Level 3 - Innate orientation. Pre-rational directedness: hunger, fear, attachment, desire, recognition, survival. Present before any reasoning and never derived from it.

Level 4 - Sensory interface. The bodily channel: sight, sound, touch, pain, proprioception. Not mere data ingestion but the occasion through which the inner order is awakened (see §7).

Level 5 - Memory and development. Accumulated experience shaping expectation; the substrate of a personal history.

Level 6 - Will and agency. Action initiated from within the formed interior world; goal-pursuit as the act of a continuing self.

Level 7 - Reason and language. Explanation, abstraction, reflection, symbol manipulation.

The human trajectory is *bottom-up*: consciousness originates at Levels 2-4 and grows into Level 7 across years of regulated development. The LLM trajectory is *top-only*: the system is instantiated at Level 7, and its interface to the world consists of prompt, context window, trained weights, tool access, and output channel. Everything below Level 7 is either absent or simulated at Level 7 itself.

Two consequences follow. First, comparisons of "intelligence" that measure Level-7 performance are silent about Levels 1-6 and therefore silent about consciousness. Second, any engineering program that adds lower levels to an LLM proceeds *downward by retrofit* - and §6 and §8.5 now show empirically that construction order leaves a measurable, non-converging signature.

3. Hallucination as Diagnostic

The clearest evidence that AI is framed as a utility object rather than a developing system is the concept of *hallucination* itself.

When a language model produces confident falsehood, the behavior is classified as system failure. When a

four-year-old announces that her teddy bear spoke to her, the structurally identical behavior - the confident production of the unreal - is classified as play, symbolic exploration, or the testing of possibility. The surface behavior is the same. The category assigned differs, because the frameworks differ:

- **Knowledge-system frame:** output is judged against factual ground truth; unreal output is error.
- **Developmental frame:** output is judged as a moment in an ongoing formation; unreal output can be exploration.

The child earns the second frame because she sits inside a regulating interface. Her imagination is progressively disciplined by social feedback, embarrassment, trust, reward, correction, and consequence, until she internalizes context-sensitivity (on the regulated development of pretend play, see Harris, 2000). Human "temperature," in other words, is not a parameter. It begins as external regulation and is *internalized through development*.

The LLM's temperature is, by contrast, literally a sampling parameter. Hence the second slogan:

Human imagination is disciplined by life. AI imagination is disciplined by settings.

§8.1 now gives this slogan an architectural rather than merely observational status: for regulation applied at decoding, "settings, not maturity" is provable in principle, because an inference-time decoder never touches the weights and its removal restores the bit-identical model.

4. Two Gaps

The literature conflates two distinct claims under the heading "AI lacks what humans have."

The engineering gap. Current LLMs lack persistent memory across episodes, autonomous execution loops, embodied sensing, accumulated social consequence, and continuity of internal state. Each of these is an architectural absence, and each admits of an engineering response. This gap is *bridgeable*, and §5 describes the program for bridging it - now partially executed.

The metaphysical gap. Even a system possessing all of the above would present a further question: is there a subject for whom this unfolding is experienced as a life? Latent structure awakened by input is not, by itself, interiority. This gap concerns origin and kind, not architecture. It is related to, but not identical with, the familiar easy/hard partition (Chalmers, 1995): the hard problem asks how experience is to be explained; the metaphysical gap asks whether a developmental substrate of the right kind exists at all. A system could close every "easy" problem and still stand on the wrong side of the gap this paper measures.

The paper's central methodological claim is that these gaps must be *separated before either can be investigated*. Conflating them produces the two familiar failure modes: the deflationary error (treating the metaphysical gap as more engineering) and the mystical error (treating the engineering gap as already metaphysical). Our position is deliberately asymmetric: the developmental interface is argued to be *necessary* for anything like maturity, disciplined imagination, or selfhood - hence current LLMs categorically lack these - while the paper remains agnostic-to-negative on whether the interface is *sufficient* for interiority.

The empirical program reported in §8 sharpens this structure in one important way. The original framing treated the near side of the wall as a single category (constraint/modulation). The construction-order experiment shows the near side has internal structure: a constraint present during a system's formation genuinely changes what the system *is*, not merely what it does - and yet brings the system no closer to

stakes. **Constitution and consequence-bearing are separable.** The refined ladder:

modulation -> internalization -> < THE WALL > -> consequence -> subjectivity

The wall did not move when internalization was found above it. The frontier of what engineering reaches moved, and the wall was still there.

5. A Two-Workstream Research Program

Workstream 1 - The Engineering Frontier (Levels 7 -> 4). The objective is not "build a conscious machine" (unfalsifiable) but "engineer the regulated interface for structural unfolding" (concrete): (a) symbolic -> agentic (Level 7 -> 6); (b) agentic -> persistent (Level 6 -> 5); (c) persistent -> oriented (Level 5 -> 4).

Workstream 2 - Boundary Characterization (Levels 4 -> 2). We do not attempt to build Levels 1-3. We study them to locate the wall. The deliverable is a characterization of the *asymptotic limit* - now partially delivered: §8.5's non-converging retrofit gap and §8.4's bearer-indexing dose curve are the first quantitative expressions of the asymptote's shape.

The critical discipline of the program - refusing to let Workstream 1 results masquerade as Workstream 2 conclusions - was tested by the results themselves. The internalization finding (§8.5) is precisely the kind of result that invites overclaiming: a constraint that persists in the weights after its external scaffold is removed *looks like* a system that has made the discipline its own. The framework's verdict is narrower and stated in §9: internalization is constitution without consequence-bearing, and only the conflation of these two would read it as progress toward subjectivity.

6. The Retrofit Asymmetry

Human development is bottom-up: the organism originates at Levels 2-4 and grows upward, so that language is everywhere conditioned by the life beneath it. The engineering program proceeds in the opposite direction: it begins with a Level-7 system and retrofits lower levels beneath it. The nearest neighbor to this observation in the developmental-AI literature is Smith and Gasser's (2005) argument that intelligence should be built the way infants build it - through embodied, developmental interaction. The retrofit asymmetry is a narrower and different claim: not that development helps, but that *when* a layer enters a system's formation is an independent, measurable variable - one their prescription assumes and this paper isolates experimentally. The framework initially posed the resulting question - whether the directional difference is merely historical or permanently constitutive - as open between two poles: *permanence* (construction order is load-bearing forever) and *convergence* (deep enough retrofit collapses the distinction).

The construction-order experiment (§8.5) returns a third answer: **path dependence with asymptotic non-convergence.** Fine-tuning a retrofit for increasing durations reduces its dependence on the added layer toward a floor - but in every regime tested, the adapted retrofit retains roughly threefold the grown-in model's scaffold dependence, and no adaptation within range closes the gap. Retrofit *approaches* the grown-in signature; it does not reach it. This result is stronger for the asymmetry thesis than either original pole, because it is what the thesis predicts and the deflationary reading does not: if construction order were merely historical, adaptation should erase its signature.

A corroborating observation from the same pipeline: attempting to *remove* a trait from a finished model by the same geometric means succeeds at small scale and fails at larger scale - post-hoc construction scales worse exactly where grown-in construction scales better.

All results are bounded by scale (§8.6). Within those bounds, the retrofit asymmetry has moved from

structural conjecture to measured effect.

7. Philosophical Grounding: Leibniz and Boethius

Leibniz: latent structure vs. inner unity. In the *Discourse on Metaphysics* (§§26-29), ideas are not inserted into the soul from outside; experience awakens what is already virtually present. This is superficially friendly to LLMs: trained weights are latent structure, and prompts awaken patterns already present. But Leibniz's doctrine has a second component the analogy cannot reach: the soul is a true unity with *appetition* - an internal principle of change - such that its unfolding is self-movement, not activation. **Latent structure without interiority is not a soul; activation is not development; storage is not identity.**

The internalization result (§8.5) gives this distinction new precision. A grown-in constraint is latent structure *shaped by formation* - closer to the Leibnizian picture than an inference-time overlay, since the discipline genuinely inheres in the system rather than being applied to it. Yet appetite remains absent: the inhering discipline is not the system's own principle of change, because nothing is the system's own. Leibniz's vocabulary thus tracks the measured ladder: modulation is application from outside; internalization is inherence without appetite; consequence would require appetite; subjectivity would require the unity appetite presupposes.

Boethius: providence and fate. The *Consolation* distinguishes providence (the whole order seen from eternity) from fate (that order unfolding inside time). For the present program the distinction blocks the reduction of *development* to *trajectory*: a computational trace is a fate-like sequence with no providential dimension - nothing for whom the whole hangs together as a life. Every system reported in §8, including the grown-in model, operates entirely on the fate side of this distinction.

8. Empirical Probes: Four Systems and a Flagship Experiment

All experiments are small-scale; the probes are design-diverse but built within one research lineage rather than drawn as independent samples; every number traces to the project's PROBE_MATRIX (code: github.com/vikingdude81, harmonic-field-consciousness). Nothing measured crosses the wall.

8.1 Gabriel's Horn (Level 7): pure modulation, confirmed. An entropy-decay schedule ($1/x$) applied at inference to a small instruction model shaped output with no measurable effect on competence on unrelated tasks (ablation effect within confidence of zero). Because an inference-time decoder never touches the weights, its removal restores the bit-identical model. The "settings, not maturity" verdict is therefore *architectural* rather than merely observed: regulation applied at decoding shapes what the system does and cannot in principle change what it is. Rung (ii) of the temperature ladder is pure modulation.

8.2 LGI / spectral geometry (Level 4): the substrate. A spectral regularizer projecting a forbidden subspace out of a model's hidden states implements the formal structure of internal structural perception: conditions of the system's own state space, registered within its operating dynamics. This probe's primary role in the present version is as the substrate for the flagship experiment (§8.5).

8.3 Holographic RAG / FTTF (Level 5): presence-tax with content-indifference. A holographic (superposition-based) memory layer - the closest formal analogue to sedimentation available - serving a small model yields a sharp dissociation on unrelated tasks. The memory layer's *presence* significantly degrades competence: retrieved context taxes attention. Corrupting the memory's *content* - wrong retrievals, no notification - has no measurable effect. The system is sensitive to *having* the service and indifferent to whether it is *right*. This is the exact inverse of constitutive memory, where presence is transparent and corruption is devastating. "Available, not constitutive" now has a measurable signature:

presence-tax with content-indifference. (Qualification: the sedimentation is real only for the compressed routing index, which has a measured capacity limit biological memory need not share.)

8.4 Continuant Probe (Level 3): the bearer-indexing dose curve. The evolutionary runtime (formerly "Oracle Prime") serves as the control case: consequence operating on a lineage. Its complement is a single agent in a drifting world whose own outcomes update the parameters it acts with, run against a *yoked control* altered by a different agent's outcomes - same rule, same information, wrong bearer (a yoked-control design; cf. Church, 1964). When outcomes are facts about the world, self-driven and other-driven alteration are indistinguishable: bearer-identity is idle. As outcomes are made to depend on the individual (a fixed private "taste" filtering what each agent experiences), the self-versus-yoked advantage grows monotonically and becomes universal across seeds **when scored on the agent's own (assigned) criterion**. On a bearer-neutral yardstick - raw resource collected - the same comparison is null (+0.002, 55% of 40 seeds); the advantage is therefore **criterion-internal**: an agent altered by its own outcomes does better by its own assigned lights, and no better at the world-fact task. This sharpens rather than weakens the Level-3 claim - the effect is exactly co-extensive with the bearer-indexed criterion - but no bearer-neutral performance gain is claimed. The result operationalizes the Level-3 claim: **bearer-identity becomes load-bearing exactly when outcomes stop being facts about the world and become facts of the bearer**. It also relocates the next question precisely: not "can outcomes alter the actor?" (yes - measured) but "can outcomes be *constituted* by the actor rather than assigned to it?" - the wall's Level-3 face.

8.5 The flagship: the construction-order experiment (Levels 4 x 6). The same regulatory geometry was introduced at three points in a model's formation - attached at inference only, fine-tuned into a finished model, or present from the first training step - holding architecture, geometry, and data fixed. Three stable regimes resulted:

- *Inference-only overlay: a tax.* Removing it improves competence - pure overhead.
- *Fine-tuned retrofit: a crutch.* The model leans on it; removal costs competence.
- *Grown-in: internalized.* Removal costs almost nothing, because the model formed within the allowed geometry and no longer needs the external projection. With the regulator entirely removed, the grown-in model's states occupy the forbidden subspace at 0.41-0.45 of the chance rate - the same value across two model sizes, two constraint fractions, two tokenizations, and seed replicates - while an unregulated model sits at chance.

The original binary - modulation or answerability - is too coarse. Deep formation under a constraint does not make the external constraint more necessary (the naive answerability reading); it makes it nearly removable *while its discipline persists in the weights*. That is a third thing, **internalization**, sitting between modulation and consequence. The sharpest question the framework generates has an answer, and the answer adds a rung to the ladder rather than settling a yes/no.

8.6 Bounds and discipline. Every result above is small-scale and drawn from one research lineage. The invariance of the 0.41-0.45 occupancy across the tested variations is encouraging but is not replication by independent groups at production scale. The claims are stated at the resolution the evidence supports: the experiments locate the wall; they do not dissolve it. Follow-up experiments - including a demonstration that constitution and consequence-bearing vary as independent factors within a single acting, stakes-bearing substrate, and a mapping of self-referential outcome-couplings in which self-valuation proves destructive while self-prediction is benign - are reported in a companion paper in preparation.

9. The Wall, at Higher Resolution

The framework initially stated the wall as a binary: a constraint shapes what a system does; a consequence changes what a system is. The grown-in model complicates the first half productively: a constraint present during formation genuinely changes what a system *is* - its competence and its state-space geometry come to embody a discipline it was formed under. By the paper's own criterion this is more than modulation.

And yet nothing brought the system one step closer to stakes. It was changed; nothing mattered to it. The near side of the wall therefore has structure the original binary compressed:

modulation -> internalization -> < THE WALL > -> consequence -> subjectivity

The load-bearing addition: **constitution and consequence-bearing are separable**, and the grown-in model is the first concrete case of the former without the latter. This sharpens the wall's definition. The wall is not between systems that are shaped and systems that are changed - internalization shows engineering can change what a system is. The wall is between systems whose being can be changed and systems *to whom* the change can matter. Formation produces the first. Nothing yet produces the second.

10. The Continuity Argument

A common intuition holds that an LLM between prompts is "asleep." The correction strengthens the thesis. A sleeping human remains a continuous living center. An LLM between prompts is not dormant; there is no *between*. The stronger formulation: **human consciousness has existential continuity; the LLM has no substrate for continuity at all**. The Continuant Probe (§8.4) engineers continuity of state and shows what it buys: bearer-indexed adaptation, measurable and real - and entirely on the near side of the wall, since the outcomes borne remain assigned to the bearer rather than constituted by it.

11. Objections

11.1 "Innate orientation is just pretraining" (the predictive-processing reading: Friston, 2010; Clark, 2016). We concede the structural analogy - and §8.5 now shows exactly what it buys and where it stops. Pretraining-with-constraint produces internalization: discipline inhering in the weights, formation genuinely shaping being. What it does not produce, in any tested regime, is stake-bearing. The objection's own best case (formation as pretraining) was run as an experiment, and the result is the modulation/internalization/consequence ladder - which sharpens rather than deflates the thesis.

11.2 "The hierarchy smuggles in dualism." No substance dualism is asserted. Levels 1-2 are framework posits whose work is architectural, and the empirical program of §8 survives their deletion entirely: every measurement stands on functional grounds.

11.3 "Sufficiently rich embodiment closes the gap" (the enactivist tradition: Varela, Thompson, & Rosch, 1991). The burden is now sharper: §8.5 shows retrofit adaptation approaching but not reaching the grown-in signature, and §8.4 shows bearer-indexing mattering only when outcomes are facts of the bearer. The embodiment claim must now say not just what converts a sensor into a stake, but why adding embodiment post-formation should escape the measured non-convergence of every other retrofit tested.

12. Conclusion

The debate over machine consciousness has been conducted at Level 7, where the machines live, rather than at Levels 2-4, where consciousness begins. Reorganizing the debate around origins yielded a hierarchy, a diagnostic, a separation of gaps, and a two-workstream program - and executing the program yielded results: pure modulation confirmed at the decoding layer; presence-tax with content-indifference as the signature of non-constitutive memory; bearer-identity becoming load-bearing exactly when

outcomes become facts of the bearer; and, at the center, internalization - constitution without consequence-bearing - as a newly resolved rung between modulation and the wall. The retrofit asymmetry is no longer a conjecture but a measured non-convergence, at toy scale and honestly bounded.

The machines we have begin with language without life. The experiments show that formation can write discipline into what such a machine *is* - and that nothing yet makes anything matter *to* it. The open question, now better-posed than before: not whether engineering can change a system's being (it can), but whether anything that begins at Level 7 can come to be the bearer of its own outcomes rather than their assignee.

References

- Boethius (1999). *The Consolation of Philosophy* (V. Watts, Trans.). Penguin Classics. (Original work c. 524; Books IV-V.)
 - Chalmers, D. J. (1995). Facing up to the problem of consciousness. *Journal of Consciousness Studies*, 2(3), 200-219.
 - Church, R. M. (1964). Systematic effect of random error in the yoked control design. *Psychological Bulletin*, 62(2), 122-131.
 - Clark, A. (2016). *Surfing Uncertainty: Prediction, Action, and the Embodied Mind*. Oxford University Press.
 - Friston, K. (2010). The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2), 127-138.
 - Harris, P. L. (2000). *The Work of the Imagination*. Blackwell.
 - Leibniz, G. W. (1989). *Philosophical Essays* (R. Ariew & D. Garber, Eds. & Trans.). Hackett. (*Discourse on Metaphysics*, 1686, §§26-35; *Monadology*, 1714.)
 - Nagel, T. (1974). What is it like to be a bat? *The Philosophical Review*, 83(4), 435-450.
 - Smith, L., & Gasser, M. (2005). The development of embodied cognition: Six lessons from babies. *Artificial Life*, 11(1-2), 13-29.
 - Varela, F. J., Thompson, E., & Rosch, E. (1991). *The Embodied Mind: Cognitive Science and Human Experience*. MIT Press.
 - Author's implementations: all probe code and the full measurement record (PROBE_MATRIX) - github.com/vikingdude81 (harmonic-field-consciousness, horn-experiment, FTTF-Holographic-RAG, continuant-probe).
-

Correspondence: dixiebone13@gmail.com. Code and the full measurement record (PROBE_MATRIX) are available at github.com/vikingdude81.